

Informed Search

Introduction to informed search
The A* search algorithm
Designing good admissible heuristics

(AIMA Chapter 3.5.1, 3.5.2, 3.6)



Outline – Informed Search

PART I - Today

- Informed = use problem-specific knowledge
- Best-first search and its variants
- A* - Optimal Search using Knowledge

- Proof of Optimality of A*
- A* for maneuvering AI agents in games
- Heuristic functions?
- How to invent them


PART II

- Local search and optimization
 - Hill climbing, local beam search, genetic algorithms,...
- Local search in continuous spaces
- Online search agents

CIS 391 - Intro to AI



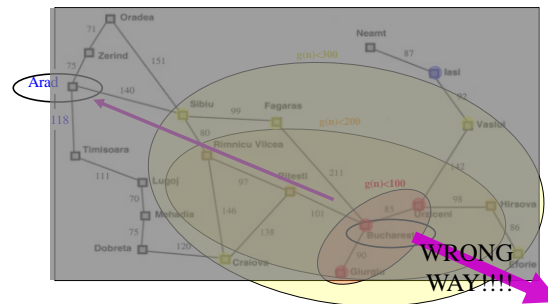
Is Uniform Cost Search the best we can do?
Consider finding a route from Bucharest to Arad..



CIS 391 - Intro to AI



Is Uniform Cost Search the best we can do?
Consider finding a route from Bucharest to Arad..



CIS 391 - Intro to AI




A Better Idea...

- Node expansion based on *an estimate* which includes distance to the goal
- General approach of informed search:
 - *Best-first search*: node selected for expansion based on an *evaluation function* $f(n)$
 - $f(n)$ includes *estimate* of distance to goal (*new idea!*)
- Implementation: Sort frontier queue by this new $f(n)$.
 - Special cases: greedy search, A* search

CIS 391 - Intro to AI



Simple, useful estimate heuristic: straight-line distances



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



Heuristic (estimate) functions

Heureka! ---Archimedes

[dictionary] "A rule of thumb, simplification, or educated guess that reduces or limits the search for solutions in domains that are difficult and poorly understood."

Heuristic knowledge is useful, but not necessarily correct.

Heuristic algorithms use heuristic knowledge to solve a problem.

A **heuristic function** $h(n)$ takes a state n and returns an estimate of the distance from n to the goal.

(graphic: <http://hyperbole.games.com/2014/10/20/eureka-moments/>)


CIS 391 - Intro to AI

7



Breadth First for Games, Robots, ...

- Pink: Starting Point
- Blue: Goal
- Teal: Scanned squares
 - Darker: Closer to starting point...



Graphics from

<http://theory.stanford.edu/~amitp/GameProgramming/>
(A great site for practical AI & game Programming)

CIS 391 - Intro to AI

9



An optimal informed search algorithm (A*)




- We add a **heuristic estimate** of distance to the goal
- Yellow: examined nodes with **high** estimated distance
- Blue: examined nodes with **low** estimated distance

CIS 391 - Intro to AI

10



Breadth first in a world with obstacles




CIS 391 - Intro to AI

11



Informed search (A*) in a world with obstacles




CIS 391 - Intro to AI

12



Greedy best-first search in a world with obstacles



CIS 391 - Intro to AI

13



Review: Best-first search

Basic idea:

- select node for expansion with minimal evaluation function $f(n)$
 - where $f(n)$ is some function that includes *estimate heuristic* $h(n)$ of the remaining distance to goal
- Implement using priority queue
- Exactly UCS with $f(n)$ replacing $g(n)$

Greedy best-first search: $f(n) = h(n)$

- Expands the node that *is estimated* to be closest to goal
- Completely ignores $g(n)$: the cost to get to n
- Here, $h(n) = h_{SLD}(n)$ = straight-line distance from to Bucharest

Greedy best-first search example

Frontier queue:
Arad 366




- Initial State = Arad
- Goal State = Bucharest

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy best-first search example

Frontier queue:

Sibiu 253
Timisoara 329
Zerind 374




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy best-first search example

Frontier queue:

Fagaras 176
Rimnicu Vilcea 193
Timisoara 329
Arad 366
Zerind 374
Oradea 380



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy best-first search example

Frontier queue:

Bucharest 0
Rimnicu Vilcea 193
Sibiu 253
Timisoara 329
Arad 366
Zerind 374
Oradea 380



Goal reached !!

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374


Properties of greedy best-first search

- **Optimal?**

- No!

— Found: Arad → Sibiu → Fagaras → Bucharest (450km)

— Shorter: Arad → Sibiu → Rimnicu Vilcea → Pitesti → Bucharest (418km)




CIS 391 - Intro to AI

20 Penn

Properties of greedy best-first search

- **Complete?**

- No – can get stuck in loops,
- e.g., Iasi → Neamt → Iasi → Neamt → ...



CIS 391 - Intro to AI

21 Penn

Properties of greedy best-first search

- **Complete?** No – can get stuck in loops,

- e.g., Iasi → Neamt → Iasi → Neamt → ...

- **Time?** $O(b^m)$ – worst case (like Depth First Search)

- But a good heuristic can give dramatic improvement of average cost

- **Space?** $O(b^m)$ – priority queue, so worst case: keeps all (unexpanded) nodes in memory

- **Optimal?** No

CIS 391 - Intro to AI

22 Penn

A* search

- Best-known form of best-first search.
- Key Idea: avoid expanding paths that are already expensive, but expand most promising first.
- Simple idea: $f(n) = g(n) + h(n)$
 - $g(n)$ the cost (so far) to reach the node
 - $h(n)$ estimated cost to get from the node to the goal
 - $f(n)$ estimated total cost of path through n to goal
- Implementation: Frontier queue as priority queue by increasing $f(n)$ (as expected...)

CIS 391 - Intro to AI

23 Penn

Admissible heuristics

- A heuristic $h(n)$ is **admissible** if it **never overestimates** the cost to reach the goal; i.e. it is **optimistic**

- Formally: $\forall n, n$ a node:

1. $h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost from n
2. $h(n) \geq 0$ so $h(G) = 0$ for any goal G .

- Example: $h_{SLD}(n)$ never overestimates the actual road distance

Theorem: If $h(n)$ is **admissible**, A* using Tree Search is **optimal**

CIS 391 - Intro to AI

24 Penn

A* search example

Frontier queue:

Arad 366

Arad 366

CIS 391 - Intro to AI

25 Penn

A* search example

Frontier queue:

Sibiu 393
Timisoara 447
Zerind 449



We add the three nodes we found to the Frontier queue.
We sort them according to the $g() + h()$ calculation.


CIS 391 - Intro to AI

26 Penn

A* search example

Frontier queue:

Rimnicu Vcea 413
Fagaras 415
Timisoara 447
Zerind 449
Arad 646
Oradea 671



When we expand Sibiu, we run into Arad again. Note that we've already expanded this node once; but we still add it to the Frontier queue again.


CIS 391 - Intro to AI

27 Penn

A* search example

Frontier queue:

Fagaras 415
Pitesti 417
Timisoara 447
Zerind 449
Craiova 526
Sibiu 553
Arad 646
Oradea 671



We expand Rimnicu Vcea.

CIS 391 - Intro to AI

28 Penn

A* search example

Frontier queue:

Pitesti 417
Timisoara 447
Zerind 449
Bucharest 450
Craiova 526
Sibiu 553
Sibiu 591
Arad 646
Oradea 671



When we expand Fagaras, we find Bucharest, but we're not done. The algorithm doesn't end until we "expand" the goal node – it has to be at the top of the Frontier queue.


CIS 391 - Intro to AI

29 Penn

A* search example

Frontier queue:

Bucharest 418
Timisoara 447
Zerind 449
Bucharest 450
Craiova 526
Sibiu 553
Sibiu 591
Rimnicu Vcea 607
Craiova 615
Arad 646
Oradea 671




Note that we just found a better value for Bucharest!
Now we expand this better value for Bucharest since it's at the top of the queue.
We're done and we know the value found is optimal!

CIS 391 - Intro to AI

30 Penn

Optimality of A* (intuitive)

- Lemma:** A* expands nodes on frontier in order of increasing f value
- Gradually adds " f -contours" of nodes
- Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$
- (After all, A* is just a variant of uniform-cost search....)



CIS 391 - Intro to AI

31 Penn

Optimality of A* using Tree-Search (proof idea)

- **Lemma:** A* expands nodes on frontier in order of increasing f value
- Suppose some suboptimal goal G_2 (i.e. a goal on a suboptimal path) has been generated and is in the frontier along with an optimal goal G .

Must prove: $f(G_2) > f(G)$

(Why? Because if $f(G_2) > f(n)$, then G_2 will never get to the front of the priority queue.)

Proof:

1. $g(G_2) > g(G)$ since G_2 is suboptimal
2. $f(G_2) = g(G_2)$ since $f(G_2) = g(G_2) + h(G_2)$ & $h(G_2) = 0$, since G_2 is a goal
3. $f(G) = g(G)$ similarly
4. $f(G_2) > f(G)$ from 1,2,3

Also must show that G is added to the frontier before G_2 is expanded – see AIMA for argument in the case of Graph Search

A* search, evaluation

- **Completeness: YES**
 - Since bands of increasing f are added
 - As long as b is finite
 - (guaranteeing that there aren't infinitely many nodes n with $f(n) < f(G)$)

A* search, evaluation

- **Completeness: YES**
- **Time complexity:**
 - Number of nodes expanded is still exponential in the length of the solution.

A* search, evaluation

- **Completeness: YES**
- **Time complexity: (exponential with path length)**
- **Space complexity:**
 - It keeps all generated nodes in memory
 - Hence space is the major problem not time

Proof of Lemma: Consistency

- A heuristic is **consistent** if

$$h(n) \leq c(n, a, n') + h(n')$$


- **Lemma:** If h is consistent,

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$

i.e. $f(n)$ is **nondecreasing** along any path.

Theorem: if $h(n)$ is consistent, **A* using Graph-Search is optimal**

Cost of getting from n to n' by any action a




A* search, evaluation

- **Completeness: YES**
 - **Time complexity: (exponential with path length)**
 - **Space complexity: (all nodes are stored)**
 - **Optimality: YES**
 - Cannot expand f_{n+1} until f_i is finished.
 - A* expands all nodes with $f(n) < f(G)$
 - A* expands one node with $f(n) = f(G)$
 - A* expands no nodes with $f(n) > f(G)$
- Also **optimally efficient** (not including ties)

Creating Good Heuristic Functions

AIMA 3.6

Heuristic functions




- For the 8-puzzle
 - Avg. solution cost is about 22 steps
—(branching factor ≤ 3)
 - Exhaustive search to depth 22: 3.1×10^{10} states
 - A good heuristic function can reduce the search process

Admissible heuristics

E.g., for the 8-puzzle:

- $h_{\text{opp}}(n)$ = number of out of place tiles
- $h_{\text{md}}(n)$ = total Manhattan distance (i.e., # of moves from desired location of each tile)




- $h_{\text{opp}}(S) = ?$
- $h_{\text{md}}(S) = ?$

Admissible heuristics

E.g., for the 8-puzzle:

- $h_{\text{opp}}(n)$ = number of out of place tiles
- $h_{\text{md}}(n)$ = total Manhattan distance (i.e., # of moves from desired location of each tile)



- $h_{\text{opp}}(S) = ?$ 8
- $h_{\text{md}}(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Relaxed problems

- A problem with fewer restrictions on the actions than the original is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_{\text{opp}}(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_{\text{md}}(n)$ gives the shortest solution

Defining Heuristics: $h(n)$

- Cost of an exact solution to a **relaxed** problem (fewer restrictions on operator)
- Constraints on **Full** Problem:
 - A tile can move from square A to square B if A is adjacent to B and B is blank.
 - Constraints on **relaxed** problems:
 - A tile can move from square A to square B if A is adjacent to B. (h_{md})
 - A tile can move from square A to square B if B is blank. (h_{opp})

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible)
 - then h_2 *dominates* h_1
- So h_2 is optimistic, but more accurate than h_1
 - h_2 is therefore better for search
 - Notice: h_{md} dominates h_{oop}
- Typical search costs (average number of nodes expanded):
 - $d=12$ Iterative Deepening Search = 3,644,035 nodes
 - $A(h_{oop}) = 227$ nodes
 - $A(h_{md}) = 73$ nodes
 - $d=24$ IDS = too many nodes
 - $A(h_{oop}) = 39,135$ nodes
 - $A(h_{md}) = 1,641$ nodes

CIS 391 - Intro to AI



Iterative Deepening A* and beyond

Beyond our scope:

- Iterative Deepening A*
- Recursive best first search (incorporates A* idea, despite name)
- Memory Bounded A*
- Simplified Memory Bounded A* - R&N say the best algorithm to use in practice, but not described here at all.
 - (If interested, follow reference to Russell article on Wikipedia article for SMA*)

(see 3.5.3 if you're interested in these topics)

CIS 391 - Intro to AI

